

Evaluating Manufacturing Success with Logistic Regression

September 28, 2019

1 Evaluating Semiconductor Manufacturing Success with Logistic Regression

In the following steps, a Logistic Regression model will be fit to predict failures in a semiconductor manufacturing facility. Because failures are typically not as frequent as many semiconductors are manufactured successfully, there will most likely be a class imbalance in the dataset that must be accounted for. To overcome this, SMOTE will be utilized to resample the training dataset and provide more failures to provide a more robust model to fit to the test dataset. Furthermore, because of the high volume of features used to describe the semiconductor manufacturing process, I also employed LASSO for feature selection to focus on only the features that provide the most impact to the manufacturing outcome. These selected features are also extremely useful information to understand which process levers are the most vital to the process. The company could then focus their efforts in improving these steps and subsequently will reduce failures.

2 Import data

The data must be imported along with the labels of each column. For now I've just named the features: feature1, feature2, and so on. These data were merged with the classification column (that holds the manufacturing outcome) and the date. All NA values were set to NaN for future imputation.

```
In [1]: # import semiconductor manufacturing features and label data
import pandas as pd
url_vars = "https://archive.ics.uci.edu/ml/machine-learning-databases/secom/secom.data"
names = ["feature" + str(x) for x in range(1, 591)] # name each feature sequentially
semi_vars = pd.read_csv(url_vars, sep=" ", names=names, na_values = "NaN") # read in

url_labs = "https://archive.ics.uci.edu/ml/machine-learning-databases/secom/secom_labels"
semi_labs = pd.read_csv(url_labs, sep=" ", names = ["classification", "date"], parse_dates=

In [2]: # merge data and take a look at the first 5 rows
semi = pd.merge(semi_vars, semi_labs, left_index=True, right_index=True)
semi.head()
```

```
Out[2]:   feature1  feature2  feature3  feature4  feature5  feature6  feature7  \
0    3030.93    2564.00    2187.7333    1411.1265      1.3602      100.0     97.6133
```

1	3095.78	2465.14	2230.4222	1463.6606	0.8294	100.0	102.3433
2	2932.61	2559.94	2186.4111	1698.0172	1.5102	100.0	95.4878
3	2988.72	2479.90	2199.0333	909.7926	1.3204	100.0	104.2367
4	3032.24	2502.87	2233.3667	1326.5200	1.5334	100.0	100.3967

	feature8	feature9	feature10	...	feature583	feature584	feature585	\
0	0.1242	1.5005	0.0162	...	0.5005	0.0118	0.0035	
1	0.1247	1.4966	-0.0005	...	0.5019	0.0223	0.0055	
2	0.1241	1.4436	0.0041	...	0.4958	0.0157	0.0039	
3	0.1217	1.4882	-0.0124	...	0.4990	0.0103	0.0025	
4	0.1235	1.5031	-0.0031	...	0.4800	0.4766	0.1045	

	feature586	feature587	feature588	feature589	feature590	classification	\
0	2.3630	NaN	NaN	NaN	NaN	-1	
1	4.4447	0.0096	0.0201	0.0060	208.2045	-1	
2	3.1745	0.0584	0.0484	0.0148	82.8602	1	
3	2.0544	0.0202	0.0149	0.0044	73.8432	-1	
4	99.3032	0.0202	0.0149	0.0044	73.8432	-1	

	date
0	2008-07-19 11:55:00
1	2008-07-19 12:32:00
2	2008-07-19 13:17:00
3	2008-07-19 14:43:00
4	2008-07-19 15:22:00

[5 rows x 592 columns]

3 Clean the dataset

After taking a look at the first 5 rows, we can see that the data have been merged properly and get an idea of how the dataset is structured. There are a few steps we can take right away to clean up the dataset a bit. I started by imputing the NaNs within each feature with the median of the respective feature. Furthermore, I edited the classification column to be more straightforward by changing the name of the column to outcome and setting the successful manufacturing outcomes as 0 and the failures as a 1.

```
In [3]: # replace missing values with the median of the column
semi.fillna(semi.median(), inplace=True)
# rename the classification column to outcome in order to be more representative of manufacturing
semi.rename(columns={'classification': 'outcome'}, inplace = True)
# map the values of the outcome column to more interpretable 0 for success and 1 for failure
semi['outcome'] = semi['outcome'].map({-1: 0, 1: 1})
```

4 Exploratory data analysis

Since the dataset describes the features that contribute to whether a semiconductor is successfully manufactured or not, I first wanted to see how the outcome was distributed. There is certainly a class imbalance within the manufacturing outcomes of this dataset with only about 100 failures out of over 1500 outcomes. This will be taken care of later on with SMOTE.

To get a general sense of the features in the dataset, I obtained the summary statistics and started to dig into a few features that seemed to have an interesting summary statistics. I plotted this histograms of a few features, feature 4 and 590 appear to be have a distribution that is skewed to the right and feature 6 appears to be made up of only one value for all runs. This is interesting to note because it will most likely be removed later since it is a constant that probably doesn't affect the manufacturing outcome.

I also plotted the manufacturing outcome as a time series to determine if there was a specific time where there was any clustering of the failures around specific time frames. There are so many rows of data it's a bit hard to parse through but there does seem to be quite a few failures in August and September of 2008.

```
In [4]: from matplotlib import pyplot as plt
```

```
# plot a histogram of the manufacturing success
plt.hist(semi['outcome'])
plt.title('Distribution of Manufacturing Success and Failures')
plt.ylabel('Count')
plt.xlabel('Manufacturing Outcome')
plt.show()
```

<Figure size 640x480 with 1 Axes>

```
In [5]: # get the exact number of manufacturing success and failures
semi['outcome'].value_counts()
```

```
Out[5]: 0    1463
        1     104
        Name: outcome, dtype: int64
```

```
In [6]: # get the summary statistics for the entire dataset
semi.describe()
```

```
Out[6]:
```

	feature1	feature2	feature3	feature4	feature5	\
count	1567.000000	1567.000000	1567.000000	1567.000000	1567.000000	
mean	3014.441551	2495.866110	2200.551958	1395.383474	4.171281	
std	73.480841	80.228143	29.380973	439.837330	56.103721	
min	2743.240000	2158.750000	2060.660000	0.000000	0.681500	
25%	2966.665000	2452.885000	2181.099950	1083.885800	1.017700	
50%	3011.490000	2499.405000	2201.066700	1285.214400	1.316800	
75%	3056.540000	2538.745000	2218.055500	1590.169900	1.518800	
max	3356.350000	2846.440000	2315.266700	3715.041700	1114.536600	

	feature6	feature7	feature8	feature9	feature10	...	\
count	1567.0	1567.000000	1567.000000	1567.000000	1567.000000	...	
mean	100.0	101.116476	0.121825	1.462860	-0.000842	...	
std	0.0	6.209385	0.008936	0.073849	0.015107	...	
min	100.0	82.131100	0.000000	1.191000	-0.053400	...	
25%	100.0	97.937800	0.121100	1.411250	-0.010800	...	
50%	100.0	101.512200	0.122400	1.461600	-0.001300	...	
75%	100.0	104.530000	0.123800	1.516850	0.008400	...	
max	100.0	129.252200	0.128600	1.656400	0.074900	...	

	feature582	feature583	feature584	feature585	feature586	\
count	1567.000000	1567.000000	1567.000000	1567.000000	1567.000000	
mean	82.403069	0.500096	0.015317	0.003846	3.067628	
std	56.348694	0.003403	0.017174	0.003719	3.576899	
min	0.000000	0.477800	0.006000	0.001700	1.197500	
25%	72.288900	0.497900	0.011600	0.003100	2.306500	
50%	72.288900	0.500200	0.013800	0.003600	2.757650	
75%	72.288900	0.502350	0.016500	0.004100	3.294950	
max	737.304800	0.509800	0.476600	0.104500	99.303200	

	feature587	feature588	feature589	feature590	outcome
count	1567.000000	1567.000000	1567.000000	1567.000000	1567.000000
mean	0.021458	0.016474	0.005283	99.652345	0.066369
std	0.012354	0.008805	0.002866	93.864558	0.249005
min	-0.016900	0.003200	0.001000	0.000000	0.000000
25%	0.013450	0.010600	0.003300	44.368600	0.000000
50%	0.020500	0.014800	0.004600	71.900500	0.000000
75%	0.027600	0.020300	0.006400	114.749700	0.000000
max	0.102800	0.079900	0.028600	737.304800	1.000000

[8 rows x 591 columns]

In [7]: *# plot a histogram of the distributions of a few features*

distribution of feature4

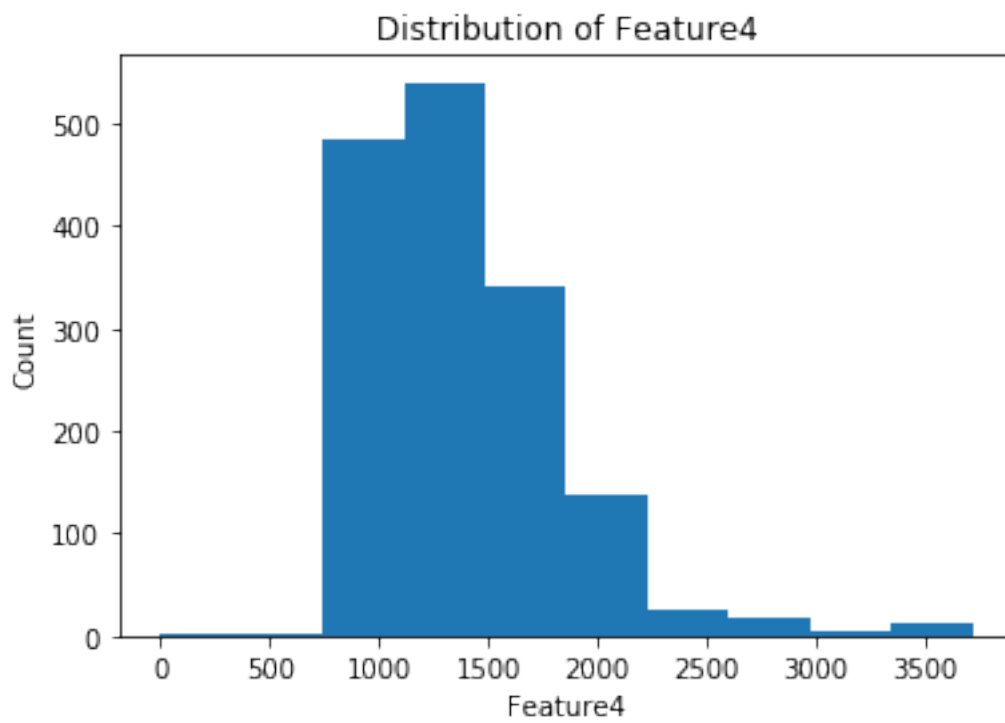
```
plt.hist(semi.feature4)
plt.title('Distribution of Feature4')
plt.ylabel('Count')
plt.xlabel('Feature4')
plt.show()
```

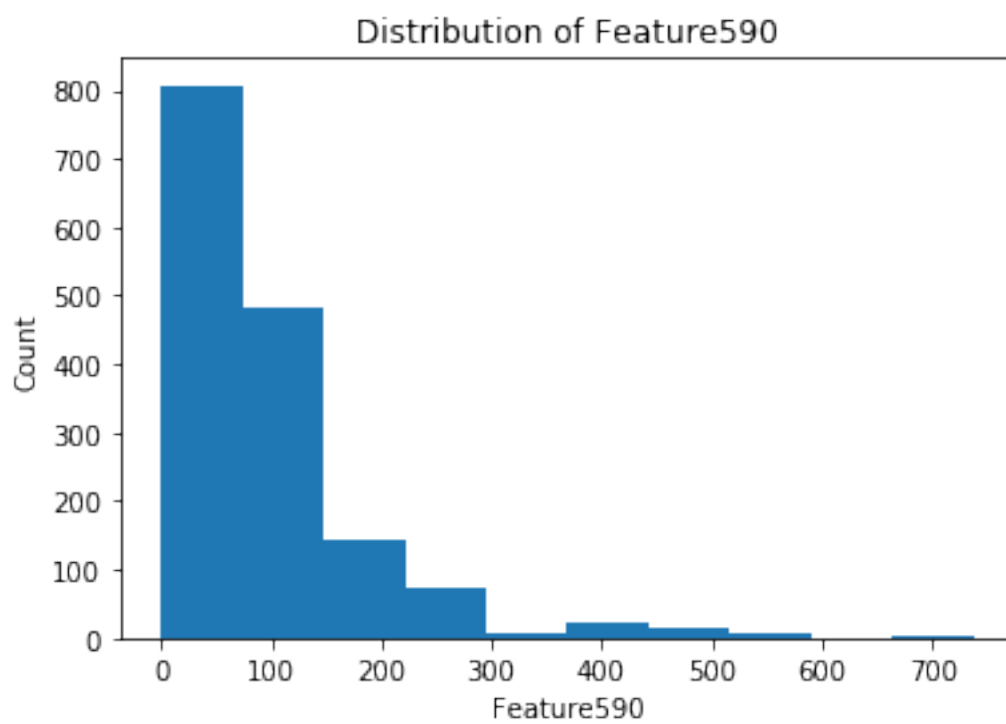
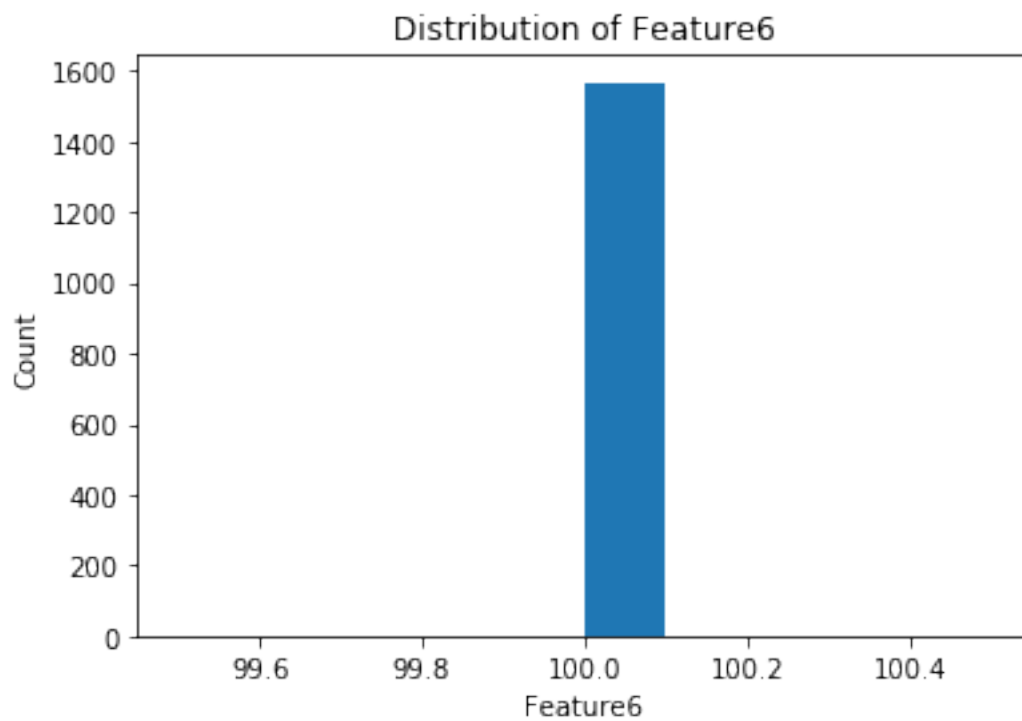
distribution of feature6

```
plt.hist(semi.feature6)
plt.title('Distribution of Feature6')
plt.ylabel('Count')
plt.xlabel('Feature6')
plt.show()
```

distribution of feature590

```
plt.hist(semi.feature590)
plt.title('Distribution of Feature590')
plt.ylabel('Count')
plt.xlabel('Feature590')
plt.show()
```





```
In [8]: # set the date column to be in the datetime format for time series analysis
        semi.loc[:, 'date'] = pd.to_datetime(semi.loc[:, 'date'])
        semi.set_index('date', inplace = True) # set index on dataset
        print(semi.head()) # print the beginning of the manufacturing dates
        print(semi.tail()) # print the end of the manufacturing dates
```

		feature1	feature2	feature3	feature4	feature5	\
date							
2008-07-19 11:55:00		3030.93	2564.00	2187.7333	1411.1265	1.3602	
2008-07-19 12:32:00		3095.78	2465.14	2230.4222	1463.6606	0.8294	
2008-07-19 13:17:00		2932.61	2559.94	2186.4111	1698.0172	1.5102	
2008-07-19 14:43:00		2988.72	2479.90	2199.0333	909.7926	1.3204	
2008-07-19 15:22:00		3032.24	2502.87	2233.3667	1326.5200	1.5334	

		feature6	feature7	feature8	feature9	feature10	...	\
date							...	
2008-07-19 11:55:00		100.0	97.6133	0.1242	1.5005	0.0162	...	
2008-07-19 12:32:00		100.0	102.3433	0.1247	1.4966	-0.0005	...	
2008-07-19 13:17:00		100.0	95.4878	0.1241	1.4436	0.0041	...	
2008-07-19 14:43:00		100.0	104.2367	0.1217	1.4882	-0.0124	...	
2008-07-19 15:22:00		100.0	100.3967	0.1235	1.5031	-0.0031	...	

		feature582	feature583	feature584	feature585	\
date						
2008-07-19 11:55:00		72.2889	0.5005	0.0118	0.0035	
2008-07-19 12:32:00		208.2045	0.5019	0.0223	0.0055	
2008-07-19 13:17:00		82.8602	0.4958	0.0157	0.0039	
2008-07-19 14:43:00		73.8432	0.4990	0.0103	0.0025	
2008-07-19 15:22:00		72.2889	0.4800	0.4766	0.1045	

		feature586	feature587	feature588	feature589	\
date						
2008-07-19 11:55:00		2.3630	0.0205	0.0148	0.0046	
2008-07-19 12:32:00		4.4447	0.0096	0.0201	0.0060	
2008-07-19 13:17:00		3.1745	0.0584	0.0484	0.0148	
2008-07-19 14:43:00		2.0544	0.0202	0.0149	0.0044	
2008-07-19 15:22:00		99.3032	0.0202	0.0149	0.0044	

		feature590	outcome
date			
2008-07-19 11:55:00		71.9005	0
2008-07-19 12:32:00		208.2045	0
2008-07-19 13:17:00		82.8602	1
2008-07-19 14:43:00		73.8432	0
2008-07-19 15:22:00		73.8432	0

[5 rows x 591 columns]

		feature1	feature2	feature3	feature4	feature5	\
--	--	----------	----------	----------	----------	----------	---

date	feature6	feature7	feature8	feature9	feature10	...	\
2008-10-16 15:13:00	2899.41	2464.36	2179.7333	3085.3781	1.4843		
2008-10-16 20:49:00	3052.31	2522.55	2198.5667	1124.6595	0.8763		
2008-10-17 05:26:00	2978.81	2379.78	2206.3000	1110.4967	0.8236		
2008-10-17 06:01:00	2894.92	2532.01	2177.0333	1183.7287	1.5726		
2008-10-17 06:07:00	2944.92	2450.76	2195.4444	2914.1792	1.5978		

date	feature582	feature583	feature584	feature585	...	\
2008-10-16 15:13:00	100.0	82.2467	0.1248	1.3424	-0.0045	...
2008-10-16 20:49:00	100.0	98.4689	0.1205	1.4333	-0.0061	...
2008-10-17 05:26:00	100.0	99.4122	0.1208	1.4616	-0.0013	...
2008-10-17 06:01:00	100.0	98.7978	0.1213	1.4622	-0.0072	...
2008-10-17 06:07:00	100.0	85.1011	0.1235	1.4616	-0.0013	...

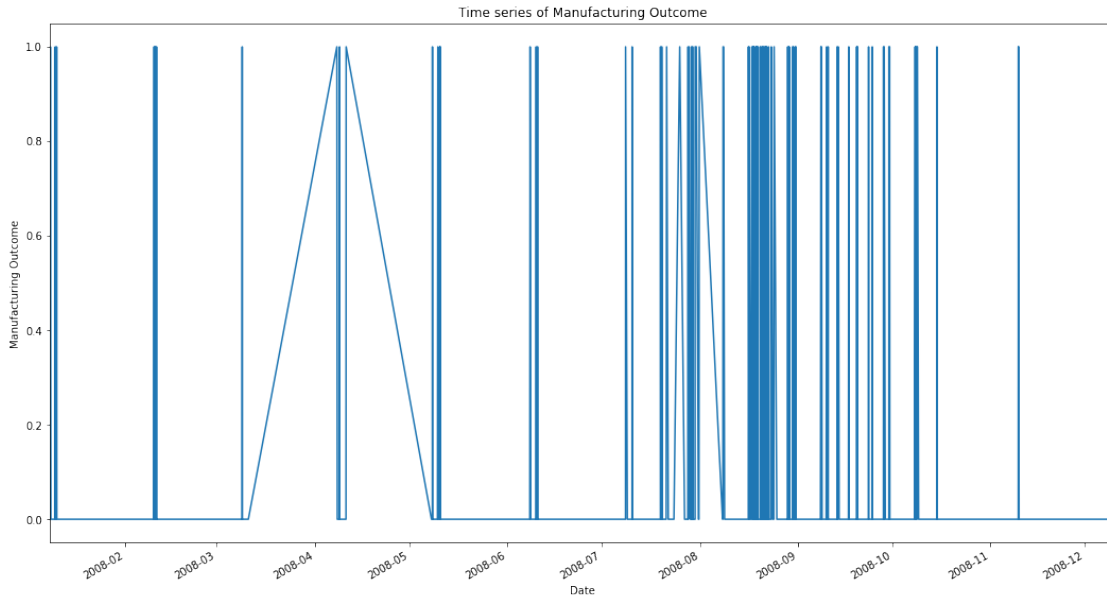
date	feature586	feature587	feature588	feature589	...	\
2008-10-16 15:13:00	203.1720	0.4988	0.0143	0.0039		
2008-10-16 20:49:00	72.2889	0.4975	0.0131	0.0036		
2008-10-17 05:26:00	43.5231	0.4987	0.0153	0.0041		
2008-10-17 06:01:00	93.4941	0.5004	0.0178	0.0038		
2008-10-17 06:07:00	137.7844	0.4987	0.0181	0.0040		

date	feature590	outcome
2008-10-16 15:13:00	203.1720	0
2008-10-16 20:49:00	203.1720	0
2008-10-17 05:26:00	43.5231	0
2008-10-17 06:01:00	93.4941	0
2008-10-17 06:07:00	137.7844	0

[5 rows x 591 columns]

```
In [9]: # plot the manufacturing outcome as a time series
ax = plt.figure(figsize=(18, 10)).gca() # define plot
semi.outcome.plot(ax = ax) # plot manufacturing outcome
ax.set_xlabel('Date')
ax.set_ylabel('Manufacturing Outcome')
ax.set_title('Time series of Manufacturing Outcome')
```


Out[9]: Text(0.5, 1.0, 'Time series of Manufacturing Outcome')



5 Prepare the data for modeling

After cleaning and getting to know the dataset a bit more, I now need to start preparing the dataset for modeling. First, I identify the outcome column as the target we are trying to predict and determine that the features will be made up of the rest of the columns in the dataset.

Next, the dataset is split up into training and test datasets further subsetting by their features and targets. The test dataset is made up of 20% of the original dataset.

```
In [10]: # establish target and features of the manufacturing data
         # set the target to the encoded manufacturing outcome column
         target = semi[['outcome']]
         # set the features as the rest of the dataset after dropping the features that are no
         feats = semi.drop(['outcome'], axis=1)

In [11]: from sklearn import model_selection

         # split original data into training and test sets
         feat_train, feat_test , target_train, target_test = model_selection.train_test_split(
                                                         test_size=0.2, random_state=6)
```

6 Perform initial Logistic Regression model

I wanted to see how the model performed without handling the class imbalance or selecting specific features. I employed a Logistic Regression model on the training data and was about 95%

accurate. This is a really good score but is most likely not very representative of the actual ability to detect failures due to the class imbalance. I will use the SMOTE method to resample the dataset and test the model performance again.

```
In [12]: from sklearn.linear_model import LogisticRegression
```

```
# instantiate logistic regression model and fit to train dataset
clf = LogisticRegression()
logR = clf.fit(feet_train, target_train)
# get the accuracy of the dataset
logR.score(feet_train, target_train)
```

```
/Users/caseythayer/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/logistic.py:433:
FutureWarning)
/Users/caseythayer/anaconda3/lib/python3.6/site-packages/sklearn/utils/validation.py:761: Data
y = column_or_1d(y, warn=True)
/Users/caseythayer/anaconda3/lib/python3.6/site-packages/sklearn/svm/base.py:931: ConvergenceW
"the number of iterations.", ConvergenceWarning)
```

```
Out[12]: 0.9497206703910615
```

7 Handling class imbalance with SMOTE

I use the SMOTE method to resample the training dataset and increase the number of failures within the dataset. This allows the model to train on more data that can help predict failures better. The resampled training data are then fed through the a Logistic Regression model again. This time the accuracy is a bit higher than before at about 96%. Because we used SMOTE, this model is a bit more representative of predicting manufacturing failures. The next step will be fine tuning the model with feature selection to see if we can improve the model further and provide the company with the features that are most vital to manufacturing failures.

```
In [13]: import imblearn
```

```
from imblearn.over_sampling import SMOTE
# handle class imbalance using SMOTE
sm = SMOTE(random_state=42)
X_res, y_res = sm.fit_sample(feet_train, target_train)
```

```
/Users/caseythayer/anaconda3/lib/python3.6/site-packages/sklearn/utils/validation.py:761: Data
y = column_or_1d(y, warn=True)
```

```
In [14]: # instantiate logistic regression model and fit to resampled training dataset
clf = LogisticRegression()
logR = clf.fit(X_res, y_res)
# get accuracy of the model
logR.score(X_res, y_res)
```

```
/Users/caseythayer/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/logistic.py:433:
FutureWarning)
/Users/caseythayer/anaconda3/lib/python3.6/site-packages/sklearn/svm/base.py:931: ConvergenceWarning:
"the number of iterations.", ConvergenceWarning)
```

```
Out[14]: 0.9597946963216424
```

8 Feature selection with LASSO

I used LASSO to select the most important features of the manufacturing dataset and used them to run another Logistic Regression model. LASSO selects the most important features by shrinking the coefficients of irrelevant features to 0 so they have little impact on the model. I still used the resampled training data obtained using SMOTE because it is more representative of data containing more failures. The model improved when using the selected features, as expected, to provide an accuracy of about 99%.

```
In [15]: # instantiate Logistic Regression model employing LASSO for feature selection
clf = LogisticRegression(penalty = 'l1')
# fit model to resampled training dataset
logR = clf.fit(X_res, y_res)
# get the accuracy of the model
logR.score(X_res, y_res)
```

```
/Users/caseythayer/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/logistic.py:433:
FutureWarning)
/Users/caseythayer/anaconda3/lib/python3.6/site-packages/sklearn/svm/base.py:931: ConvergenceWarning:
"the number of iterations.", ConvergenceWarning)
```

```
Out[15]: 0.9850299401197605
```

9 Conclusion

Predicting semiconductor manufacturing failure rates requires quite a few steps to get to a representative and robust model. In this case study, I cleaned the dataset, performed some EDA and prepared the data for modeling. Next, I evaluated a Logistic Regression model with a number of steps to improve its ability to predict manufacturing failures. First, I tested the model without changing anything to figure out what the starting point was and how the model was performing straight up (accuracy around 96%).

Through EDA, I observed a class imbalance in the manufacturing outcome. I handled the class imbalance by employing SMOTE to resample the test dataset and provide more data to train on. I tested the model again with the resampled data and found the model accuracy decreased a bit, which was expected (94%). The resampled data provides a more representative number of failures to train on so it helps train the model in a more realistic sense.

Finally, I used LASSO to select the most important features. LASSO selects the most vital features by setting the coefficients of the irrelevant features to 0 and therefore, removing their

impact on the model. After performing LASSO, I fit the model to the resampled data again and found that it improved the accuracy even more (99%). In addition, the list of most important features can be provided to the company to guide future improvements in the process and prevent failures.